

# Implementing XKCD 1425 with YOLO, Darknet and Python

Colin Sauze  
[<cos@aber.ac.uk>](mailto:<cos@aber.ac.uk>)



SUPERCOMPUTING WALES  
UWCHGYFRIFIADURA CYMRU



# Overview



IN CS, IT CAN BE HARD TO EXPLAIN  
THE DIFFERENCE BETWEEN THE EASY  
AND THE VIRTUALLY IMPOSSIBLE.

- Motivations
- Has this been done before?
- YOLO and Darknet
- Are we in a national park?
- Putting it all together
- Results

From: <https://xkcd.com/1425/>

Published: September 24<sup>th</sup> 2014  
(LESS THAN FIVE YEARS AGO!)

# Motivations

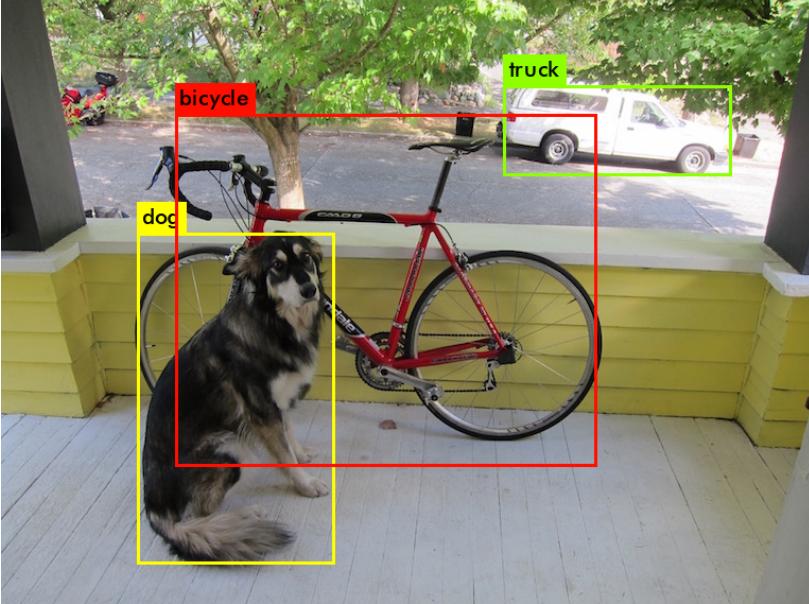
- I'm a Research Software Engineer for Super Computing Wales
- Help researchers across the university use High Performance Computing
- Lots of people asking me about Deep Learning/CNNs/Image recognition
- Decided to read up on it
  - Found Darknet/YOLO
  - Realised implementing XKCD1425 would be easy with this
  - Nice demo to show what it can do

# Has this been done before?

The screenshot shows a blog post on the code.flickr.com website. The title is "Introducing: Flickr PARK or BIRD". It features two images: one of a Zion National Park sign and another of a Secretary Bird. A large "OR" is placed between the two images. The post includes a link to "parkorbird.flickr.com!"

- October 2014, Flickr released “Park or bird”
- Tells you if an image is a park OR a bird
- Comic said “In a park AND a bird”
- **<http://parkorbird.flickr.com/>**  
No longer online

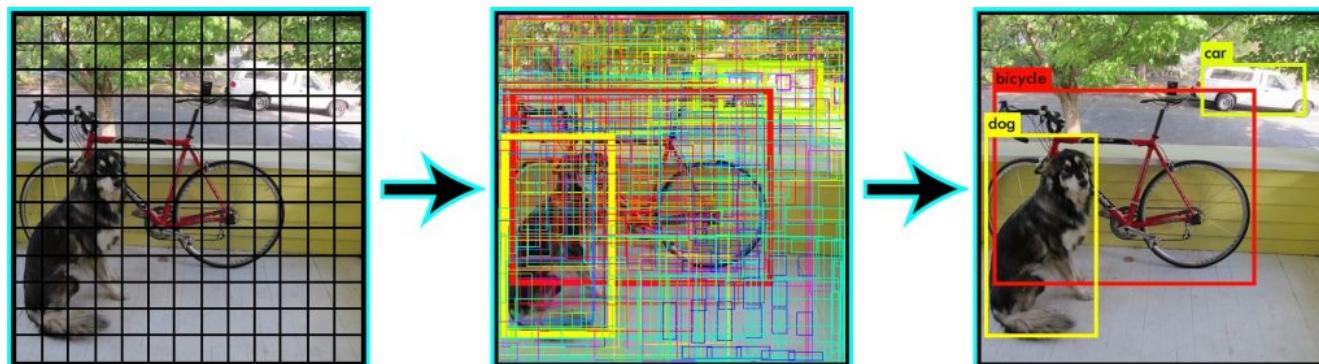
# YOLO and Darknet



- Darknet very efficient neural network library
- YOLO = You Only Look Once
  - Object recognition + localisation using Darknet
  - Single neural network does the entire task
- <https://pjreddie.com/darknet/yolo/>
- Video demo:
  - <https://www.youtube.com/watch?v=VOC3huqHrss>

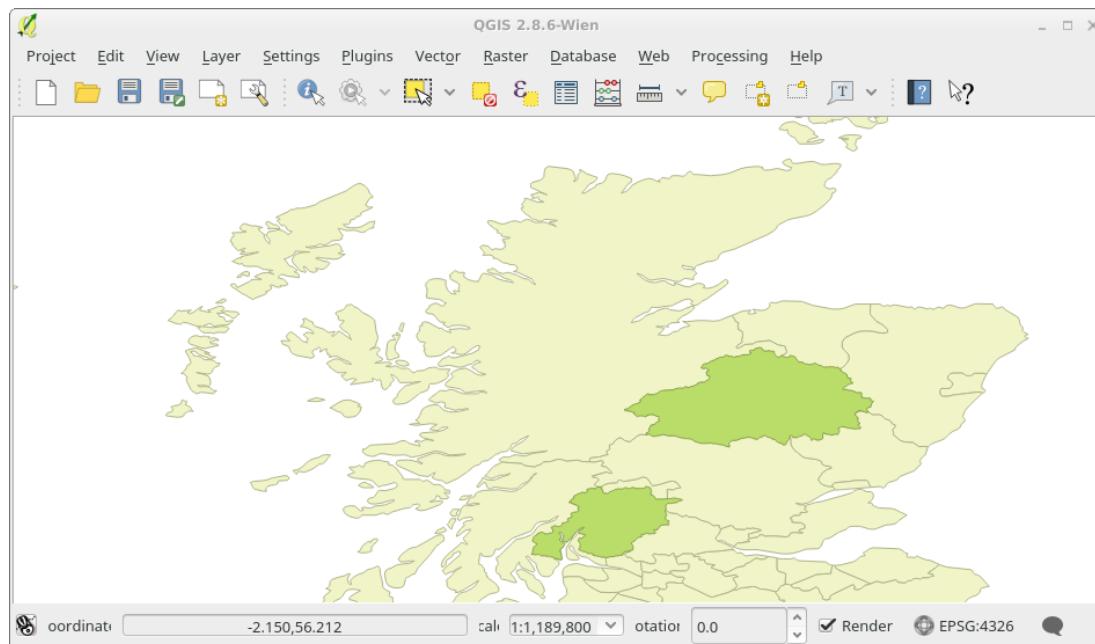
# How YOLO Works

- Single Neural Network processes whole image
- Divides image into regions, makes bounding boxes
- Predicts what class that box belongs to
- Faster than alternatives, comparable accuracy
- Trained on Common Objects in CONext (COC) dataset
  - 300k tagged, segmented images
  - 80 categories, including birds

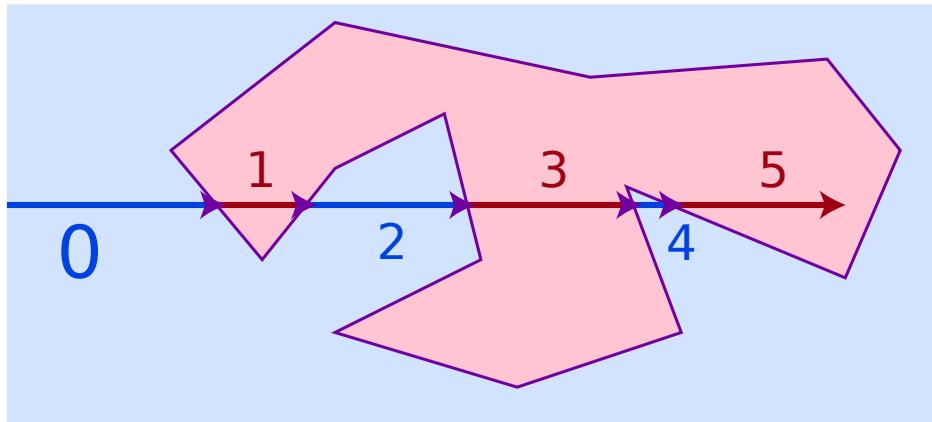


# Getting Park Boundary Data

- Data.gov.uk UK national park boundary data
  - Each park is separate download
  - Haven't looked at other countries yet
  - Open Street Map could supply data too
- Converted coordinate format with GDAL (Geospatial Data Abstraction Library)

A screenshot of a web browser displaying the data.gov.uk website. The address bar shows the URL "https://data.gov.uk/search?". The page header reads "data.gov.uk | Find open data" and "BETA". A message at the top says "This is a new service – your [feedback](#) will help us to improve it". Below this is a section titled "Search results" with a search bar containing "national park boundary". To the right of the search bar is a magnifying glass icon. Further down, there are sections for "Filter by", "Publisher", "Topic", and "Published by". The "Publisher" dropdown is set to "National Park Boundary". The "Topic" dropdown is set to "Peak District National Park Authority". At the bottom, a timestamp indicates "Last updated: 29 March 2016".

# Are we in a national park?



```
# read the shapefiles
parks[0]= shape(shapefile.Reader("Cairngorms.shp").shapes()[0])
parks[1]= shape(shapefile.Reader("LochLomond.shp").shapes()[0])
#store the park names
parknames = [ "Cairngorms", "Loch Lomond"]

def get_park_name(lat, lon):
    """Returns the name of the park this point is in"""
    point = Point(lon, lat)
    for i in range(len(parks)):
        polygon = parks[i]
        if polygon.contains(point):
            return parknames[i]
```

- Latitude/Longitude of an image from its EXIF JPEG header
  - Assumes camera has a GPS and its accurate
- Draw a line from the point through the park boundary polygon
  - Even number of intersections = outside park
  - Odd number = inside park
- Python Shapely library does this for us

# Putting it all together

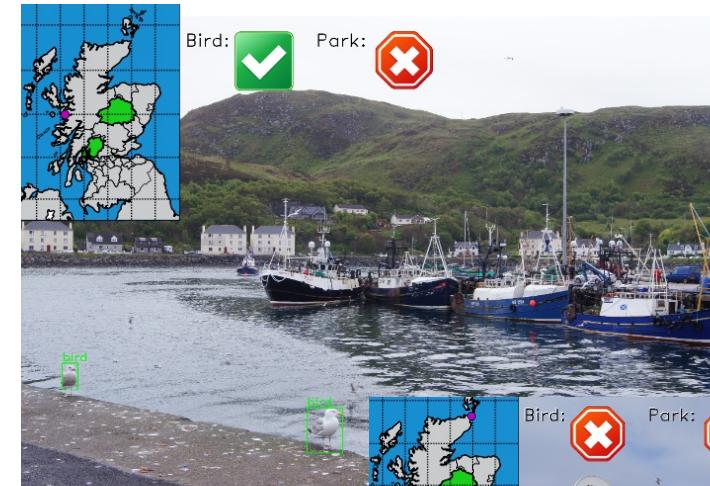
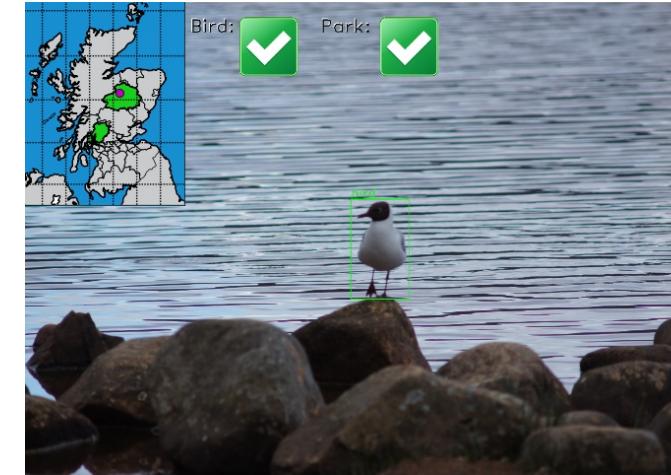
- Approximately 200 lines of Python
  - Mostly utility stuff like loading files
- Darknet code requires 4 lines:

```
dn.set_gpu(0)  
net = dn.load_net("cfg/yolov3.cfg", "yolov3.weights", 0)  
meta = dn.load_meta("cfg/coco.data")  
r = dn.detect(net, meta, "image.jpg")
```

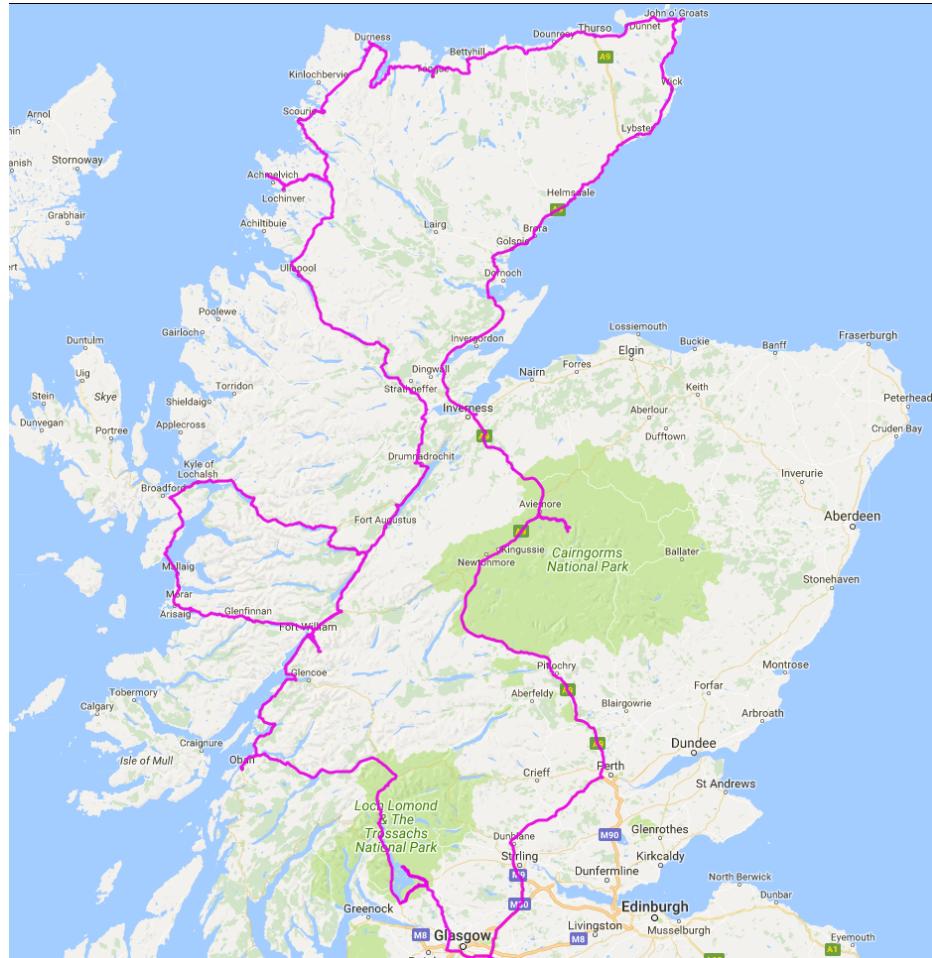
- Annotated image output
- Outputs a CSV file listing each image

## Basic algorithm:

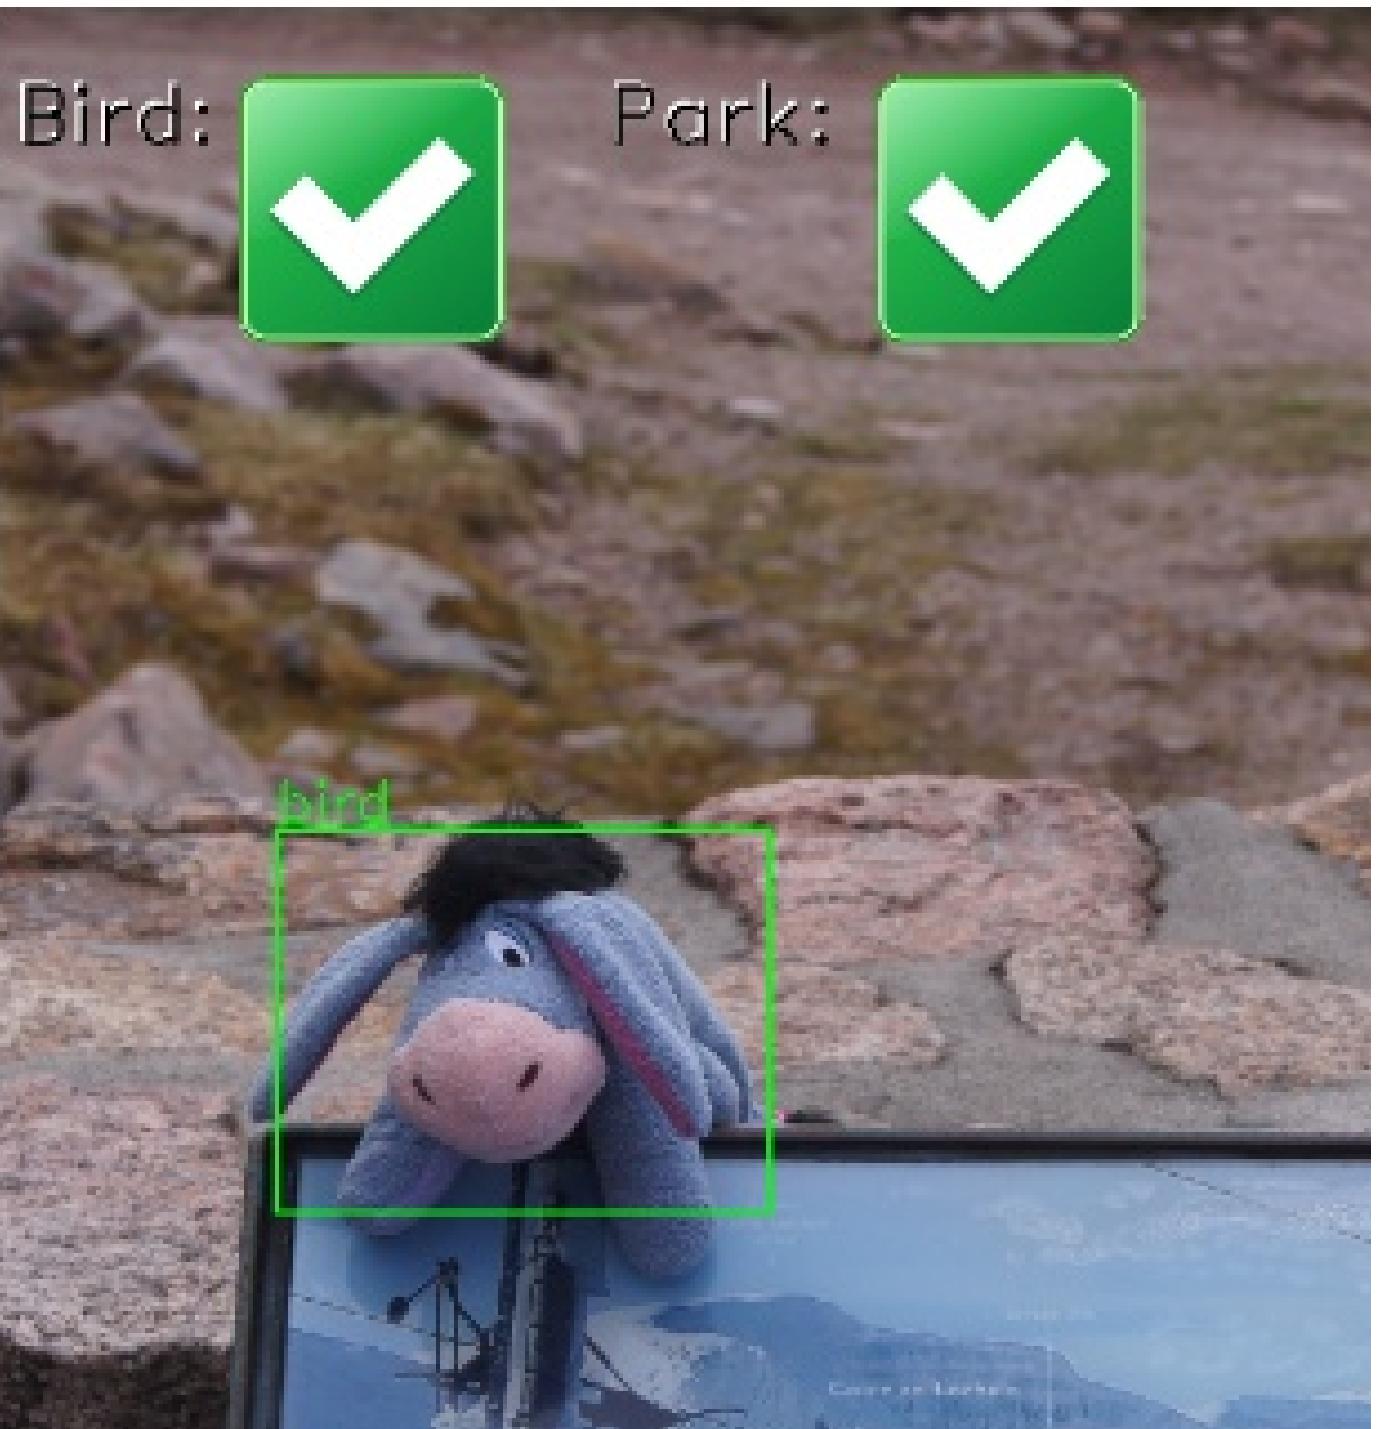
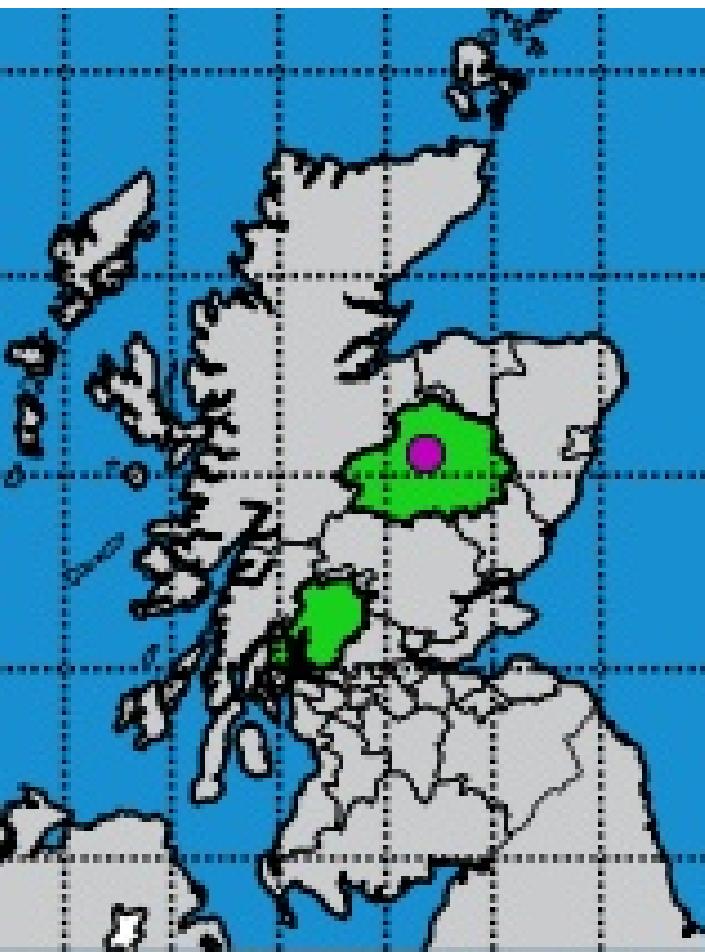
- Are we in a park?
  - Yes: Run bird detection
    - Was a bird found?
    - Yes: Draw bounding box
    - No: Go to next image
  - No: go to next image

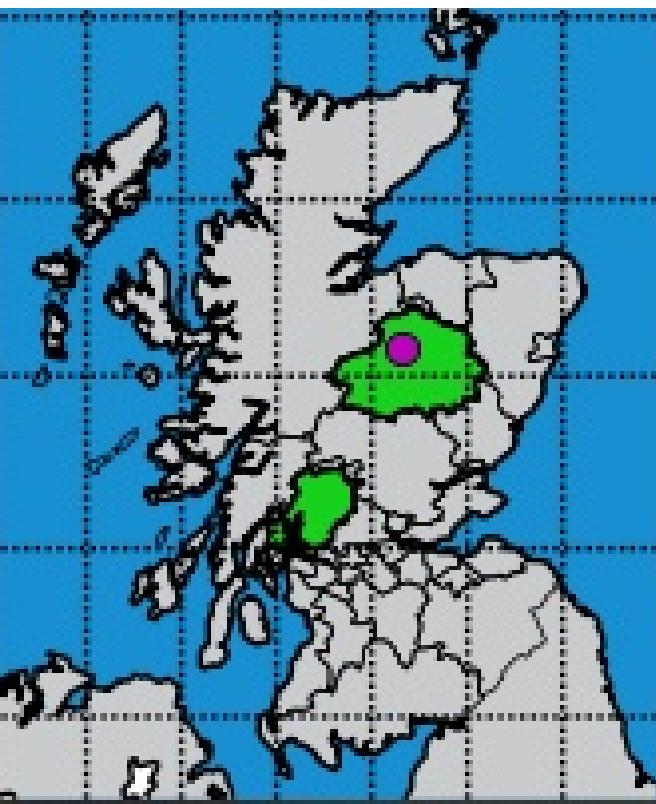


# Testing – My Scottish Holiday 2015



- One week camping in Scotland
- Visited Cairngorm and Loch Lomond National Parks
- 506 Geotagged images from Sony A55 camera
- 35 contain birds
- 12 identified correctly
- 12 false positives
- 23 false negatives

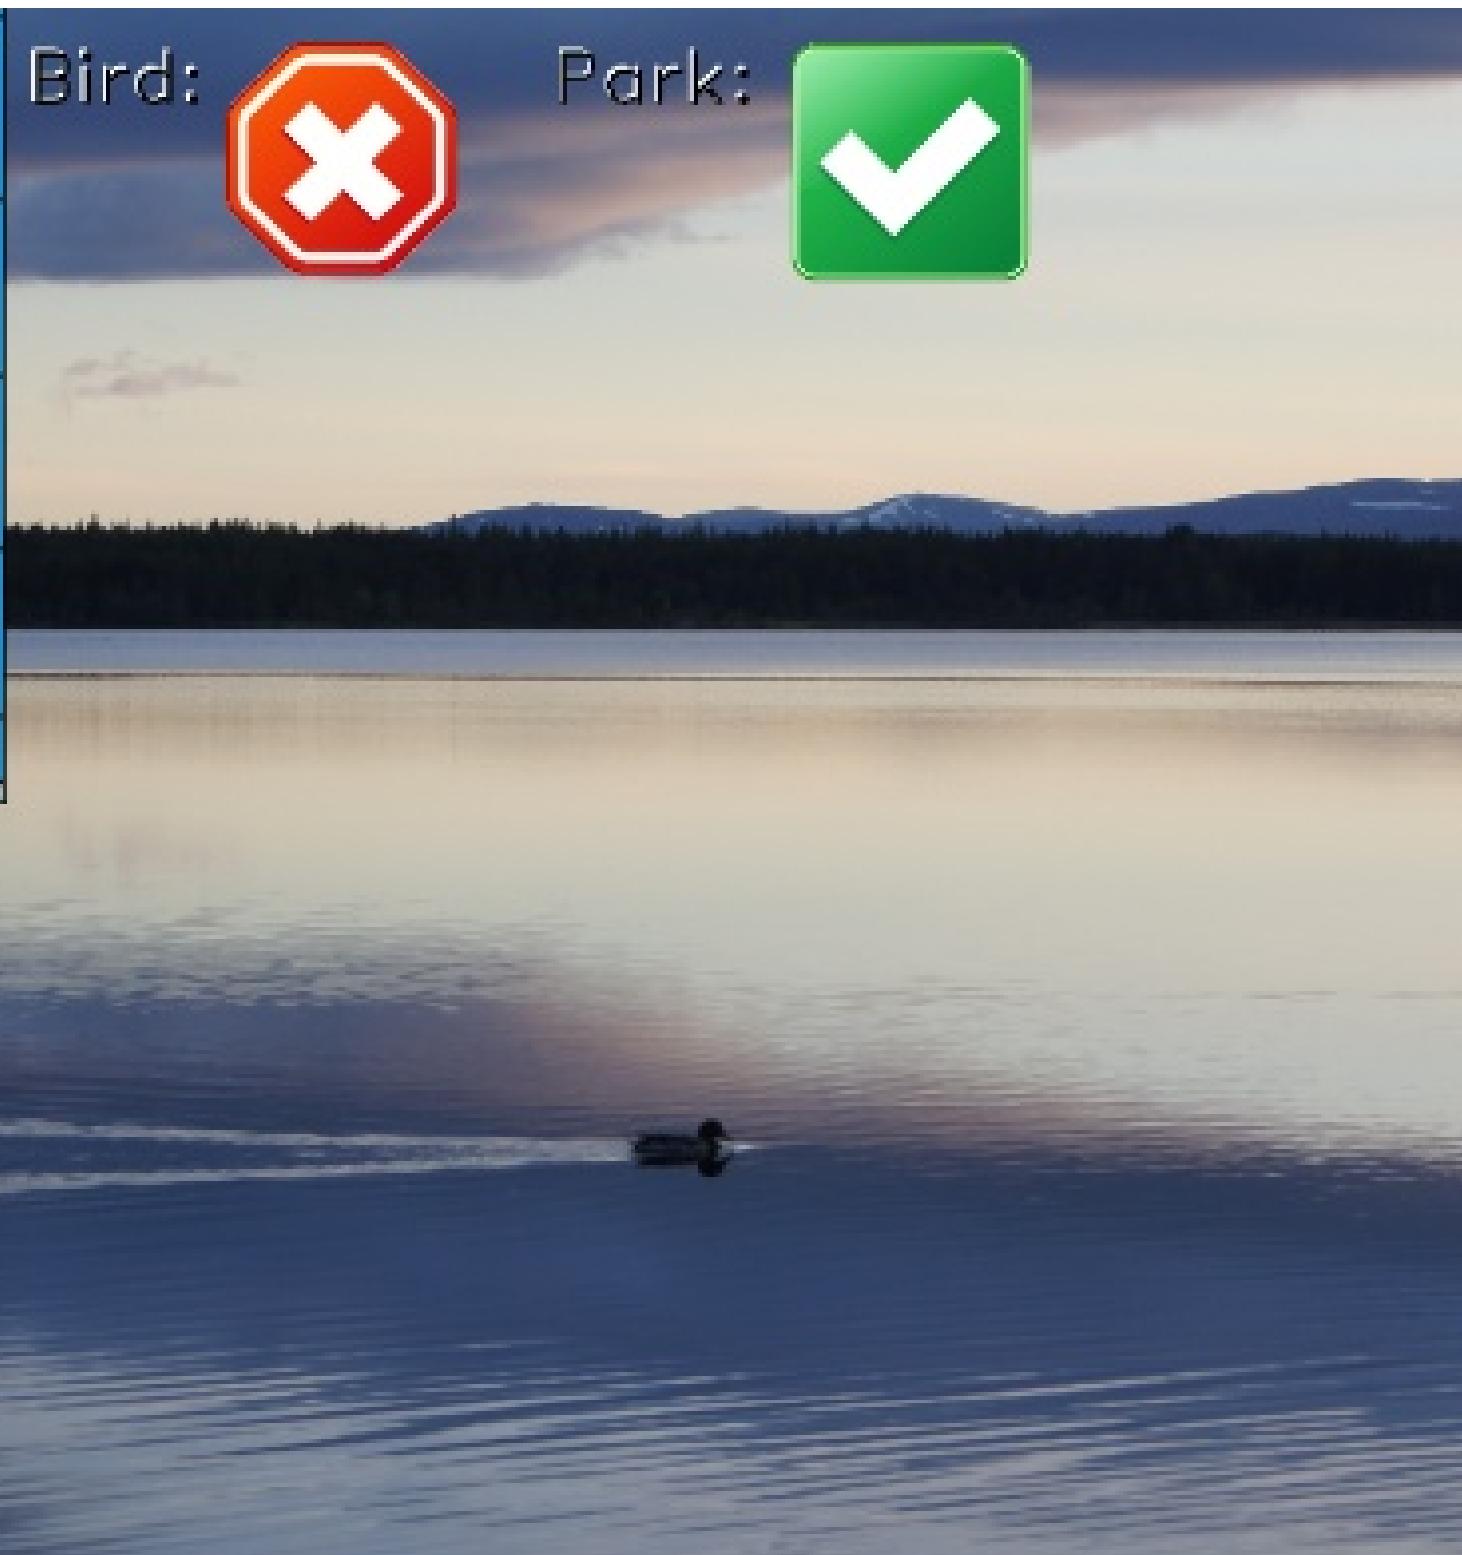


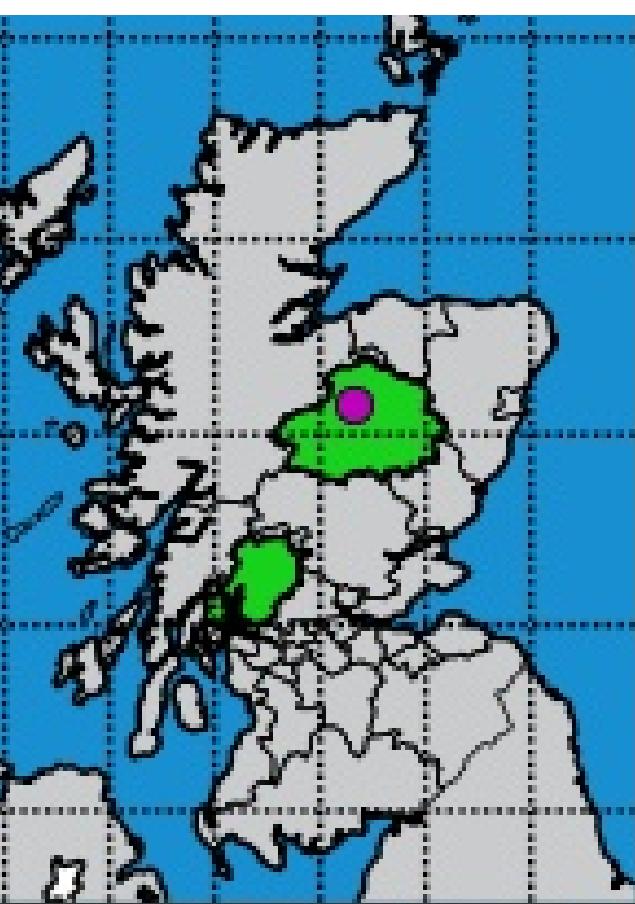


Bird:



Park:

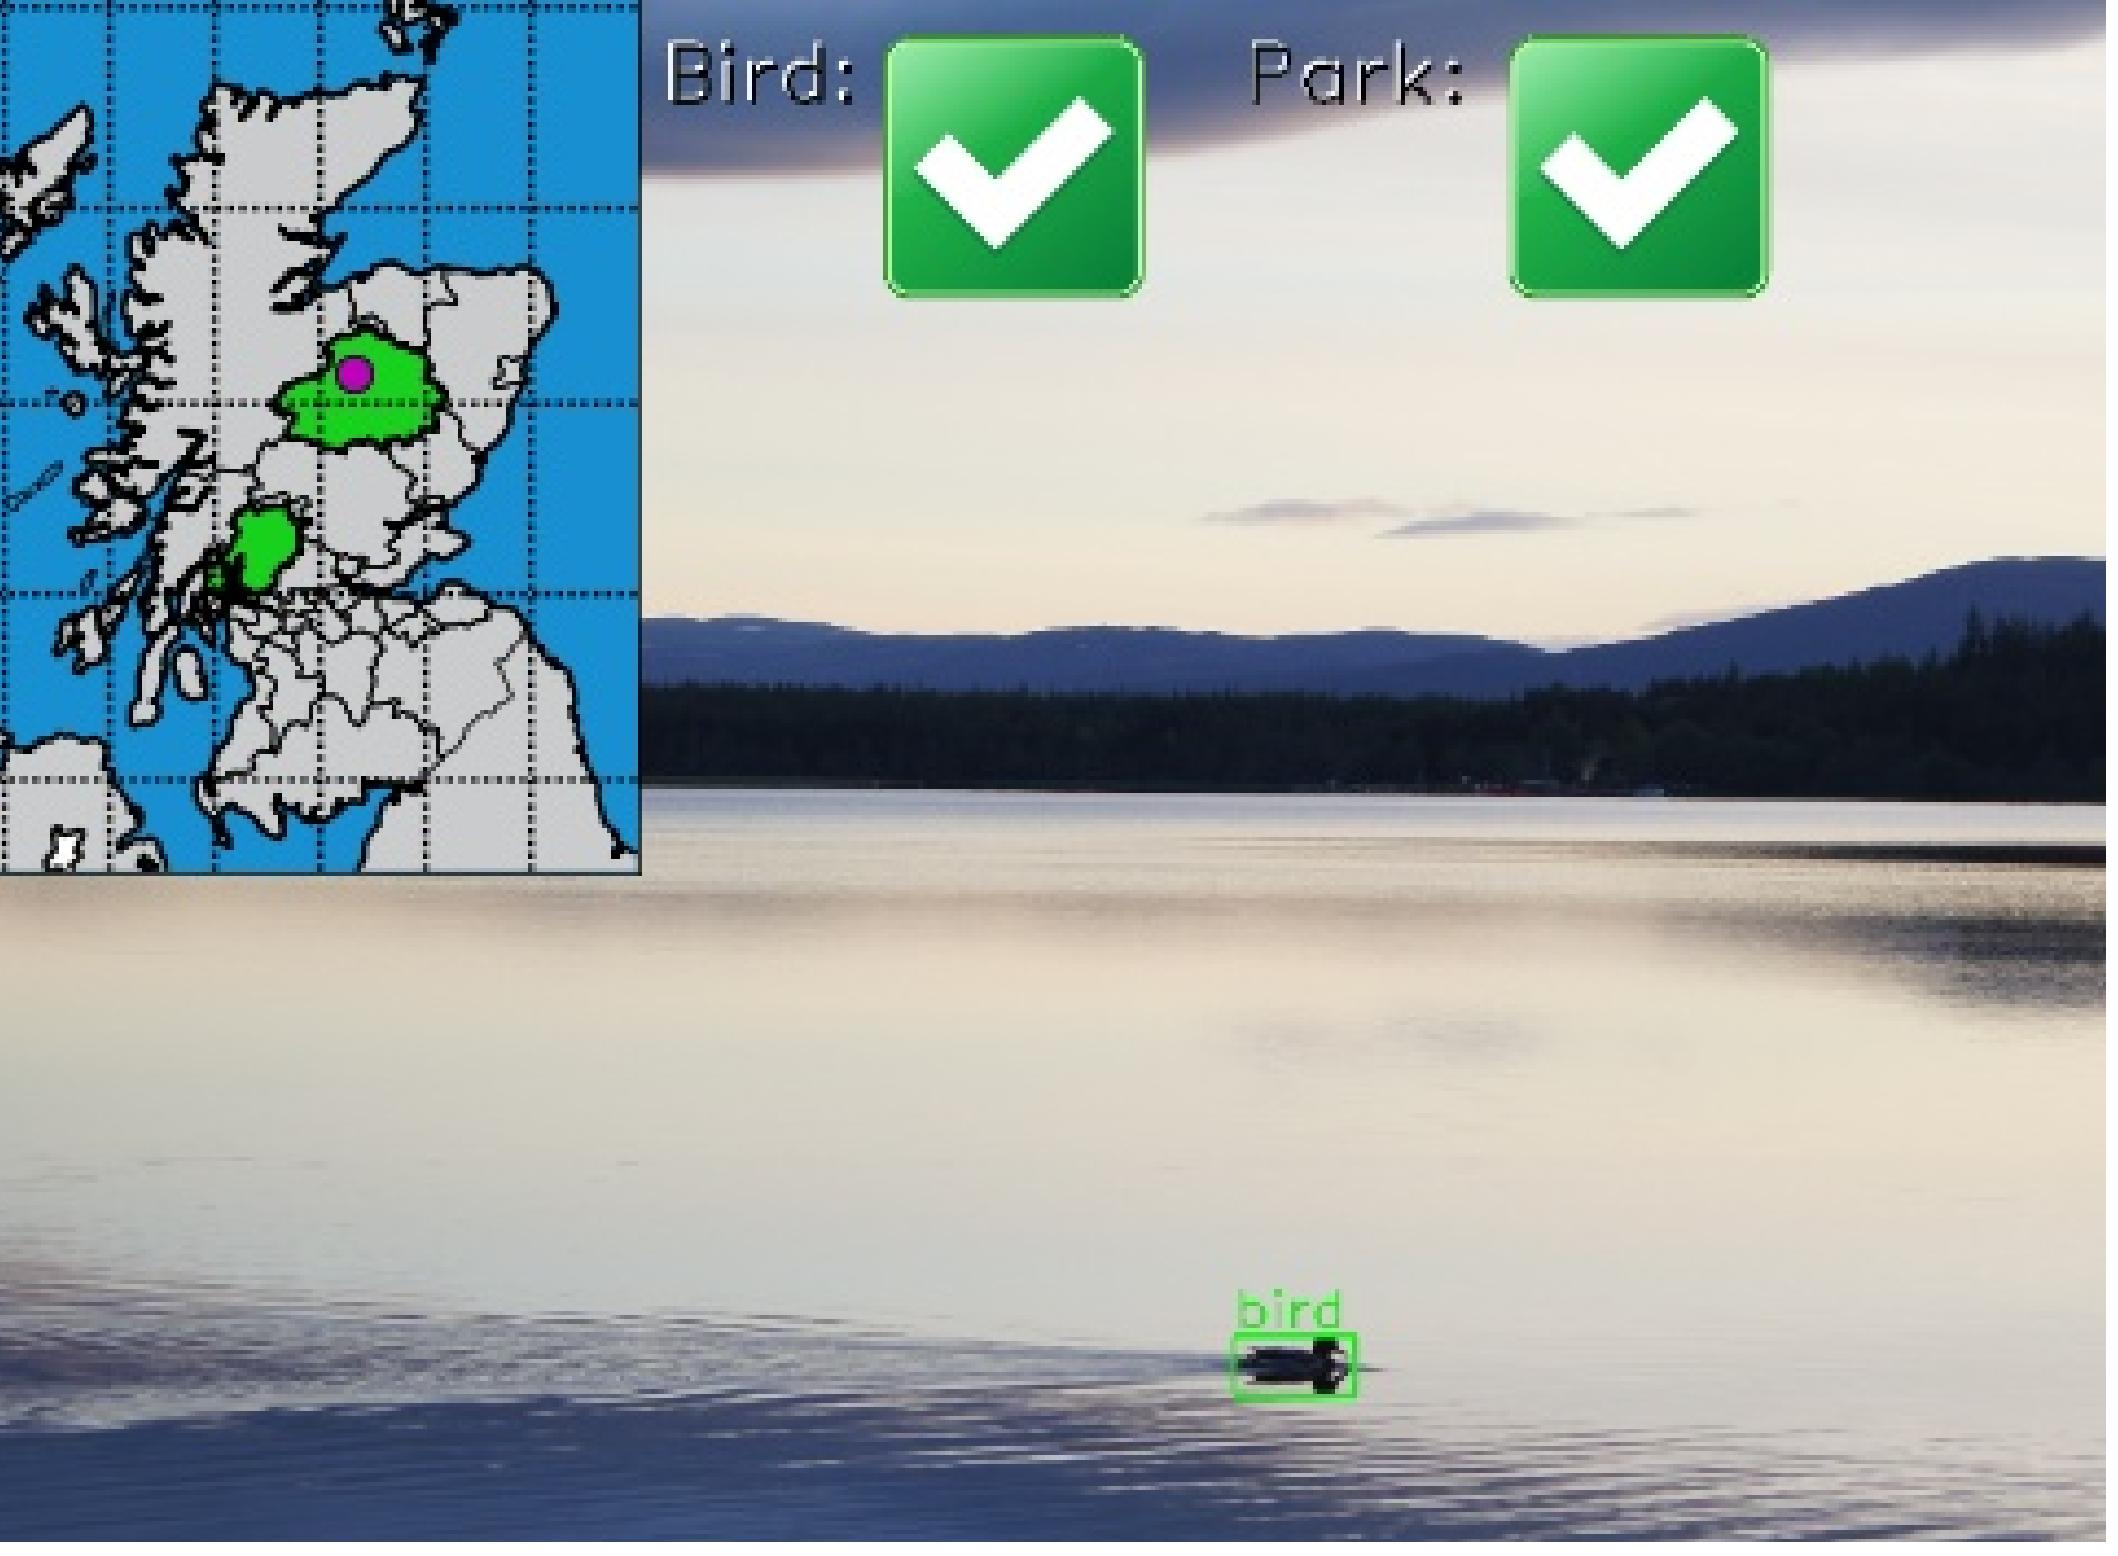


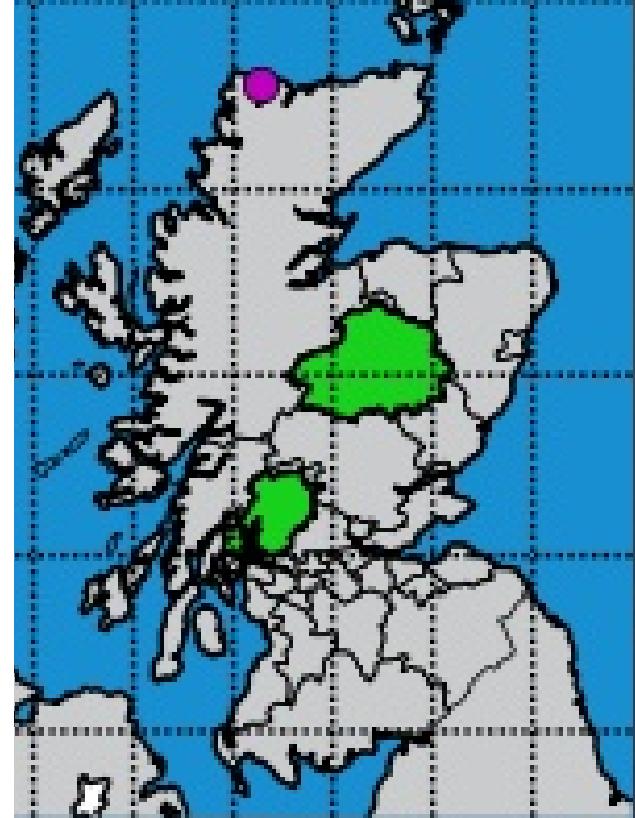


Bird:



Park:



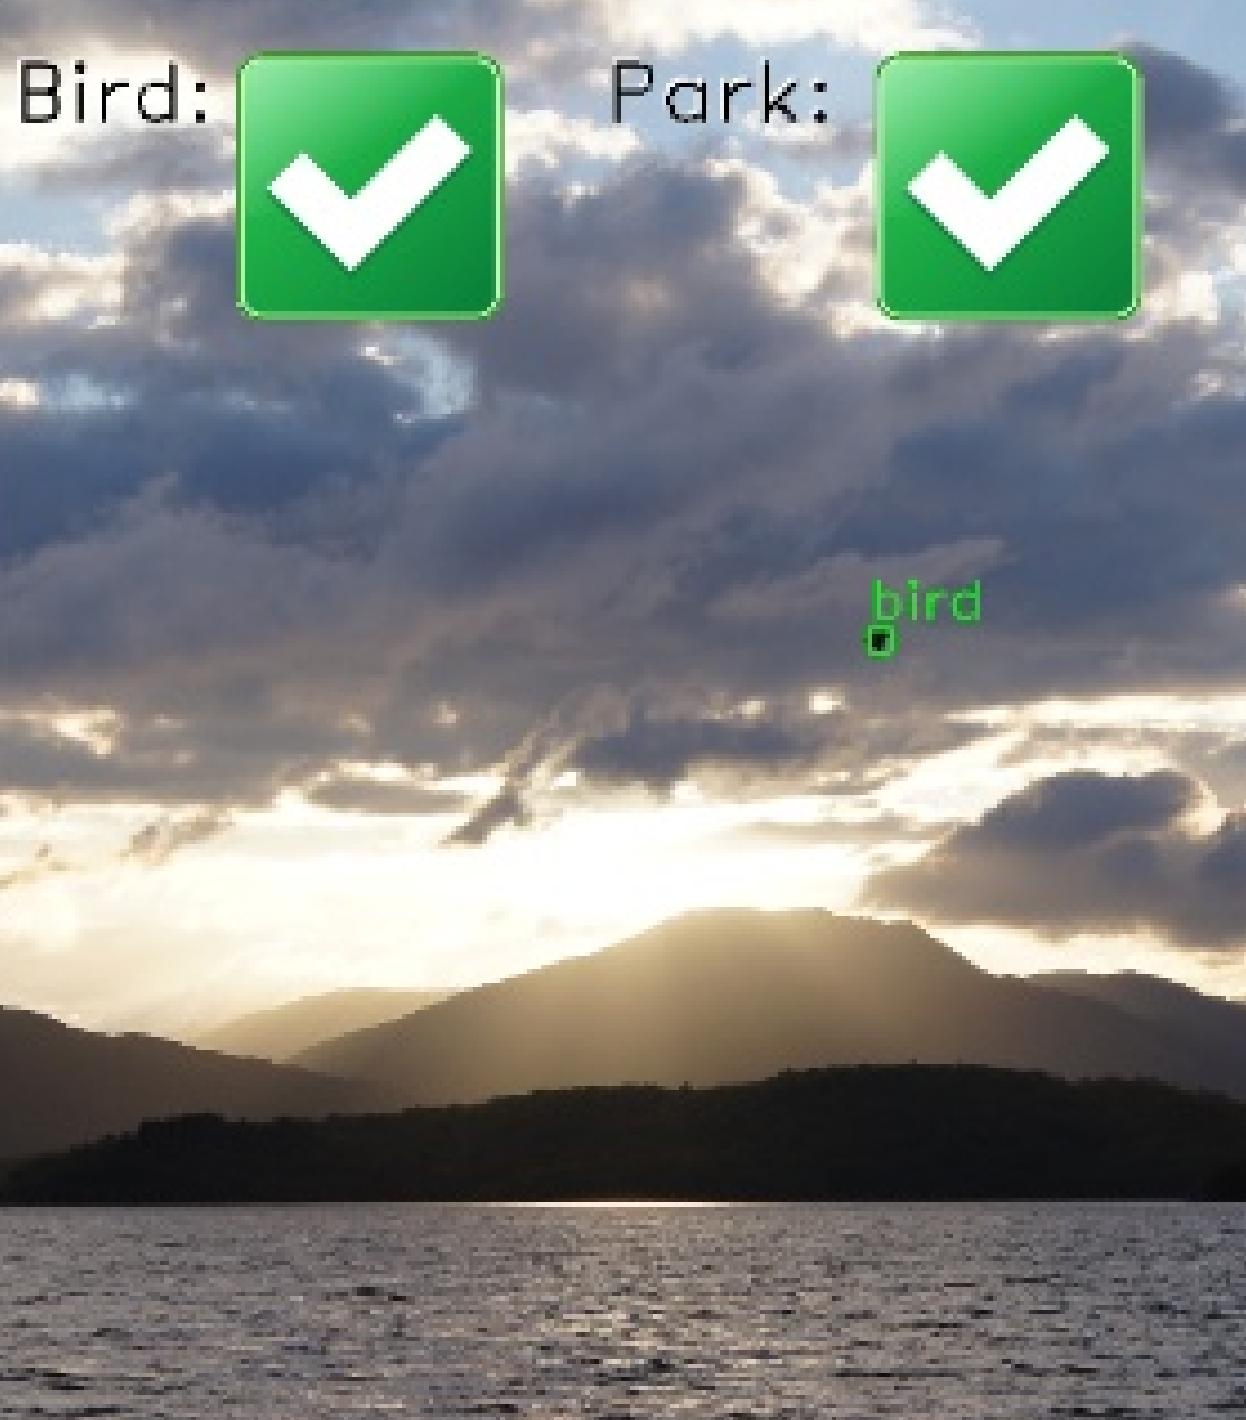
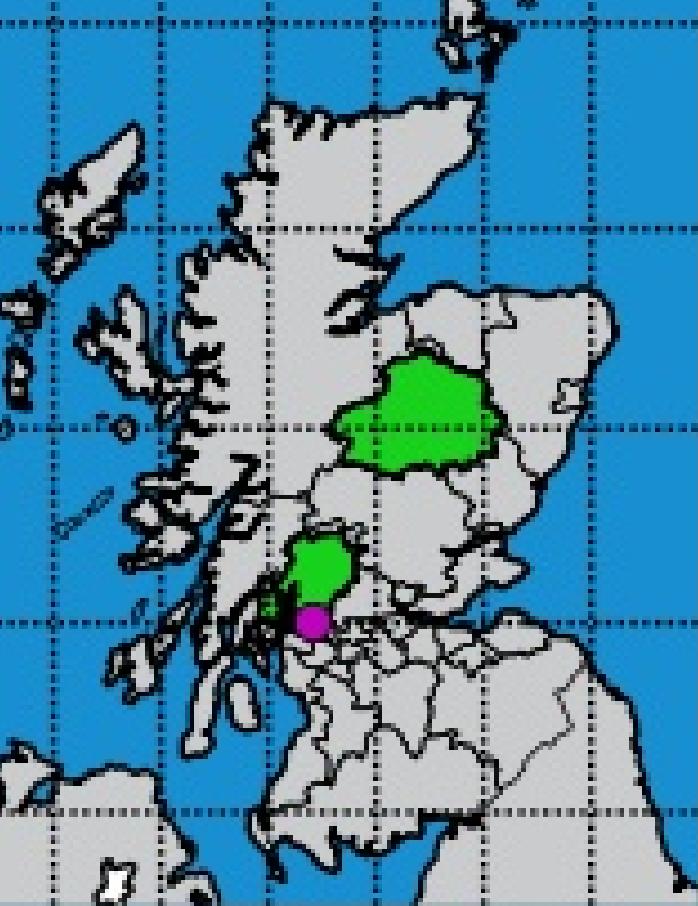


Bird:



Park:





Bird:



Park:



# Performance

- YOLO/Darknet can do real time 30fps on a GPU
  - I haven't got a suitable GPU....yet
- 2 hours, 40 minutes to check 506 images
  - Intel i7-7500U CPU, quad core but only using one core.
  - Just under 19 seconds per image

# Conclusions

- Not perfect performance
  - Some of the false negatives surprised me
  - Stuffed donkeys are apparently birds!
  - Threshold changes might have reduced false negative rate
- I'm still impressed by YOLO
  - Writing this was easy
  - I did no training of neural nets or tweaking of threshold parameters
- COCO only has 3362 bird images, most fairly close up
  - <http://cocodataset.org/#explore>
  - Could use ImageNet data.
    - Millions of images
    - Not annotated, no localisation
- Code available from <https://github.com/colinsauze/xkcd1425-with-yolo>